

Weiyan Shi

CONTACT	ws2634@columbia.edu wyshi.github.io Google Scholar	
RESEARCH INTERESTS	· Dialogue systems, especially social influence dialogue, e.g., persuasion · Safety in NLP, e.g., privacy-preserving models	
EDUCATION	Columbia University Ph.D. in Computer Science Advisor: Zhou Yu	2020 - 2023 (anticipated)
	University of California, Davis Ph.D. in Computer Science (transferred out)	2018 - 2020
	University of California, Berkeley M.A. in Statistics	2015 - 2016
	Renmin University of China B.S. in Mathematics and Applied Mathematics	2011 - 2015
AWARDS & HONORS	Best Paper Nomination, ACL Dean's Distinguished Ph.D. Fellowship Department Citation (top 1), UC, Berkeley Speaker at Department Commencement (top 1), UC, Berkeley National Scholarship, RUC Presidential Fellowship for Studying Abroad, RUC	2019 2018-2023 2016 2016 2014 2013
PRIOR POSITIONS	Research Intern , Facebook AI Research <i>Host</i> : Jason Weston and Jing Xu <i>Project</i> : Integrating user feedback into chatbots	2022.05 - 2022.09
	Research Intern , Facebook AI Research <i>Host</i> : Mike Lewis <i>Project</i> : Zero-shot dialogue nonsense detection	2021.05 - 2021.12
	Research Intern , Facebook AI Research <i>Host</i> : Mike Lewis <i>Project</i> : Strategic negotiation dialogue generation	2020.06 - 2020.12
	Full-time Data Scientist , [24]7.ai Customer service chatbot development	2016 - 2018

PUBLICATIONS *The listed conferences are top-tier in NLP with acceptance rates between 20%-25%.*

[Selective Differential Privacy for Language Modeling](#)

Weiyan Shi, Aiqi Cui, Evan Li, Ruoxi Jia, Zhou Yu

North American Chapter of the Association for Computational Linguistics (NAACL), 2022

Refine and Imitate: Reducing Repetition and Inconsistency in Persuasion Dialogues via Reinforcement Learning and Human Demonstration

Weiyan Shi, Yu Li, Saurav Sahay, Zhou Yu

Empirical Methods in Natural Language Processing (EMNLP) 2021 Findings

LEGOEval: An Open-Source Toolkit for Dialogue System Evaluation via Crowdsourcing

Yu Li, Josh Arnold, Feifan Yan, Weiyan Shi, Zhou Yu

Annual Meeting of the Association for Computational Linguistics (ACL) 2021 Demo

PRAL: A tailored pre-training model for task-oriented dialogue generation

Jing Gu, Qingyang Wu, Chongruo Wu, Weiyan Shi, Zhou Yu

Annual Meeting of the Association for Computational Linguistics (ACL) 2021

INSPIRED: Toward Sociable Recommendation Dialogue Systems

Shirley Anugrah Hayati, Dongyeop Kang, Qingxiaoyang Zhu, Weiyan Shi, Zhou Yu

Empirical Methods in Natural Language Processing (EMNLP) 2020

Structured Attention for Unsupervised Dialogue Structure Induction

Liang Qiu, Yizhou Zhao, Weiyan Shi, Yuan Liang, Feng Shi, Tao Yuan, Zhou Yu, Song-Chun Zhu

Empirical Methods in Natural Language Processing (EMNLP) 2020

Understanding User Resistance Strategies in Persuasive Conversations

Youzhi Tian, Weiyan Shi, Chen Li, Zhou Yu

Empirical Methods in Natural Language Processing (EMNLP) 2020 Findings

Effects of Persuasive Dialogues: Testing Bot Identities and Inquiry Strategies

Weiyan Shi, Xuwei Wang, Yoo Jung Oh, Jingwen Zhang, Saurav Sahay, Zhou Yu

Conference on Human Factors in Computing Systems (CHI) 2020

End-to-End Trainable Non-Collaborative Dialogue System

Yu Li, Kun Qian, Weiyan Shi, Zhou Yu

AAAI Conference on Artificial Intelligence (AAAI) 2020

How to Build User Simulators to Train RL-based Dialogue Systems

Weiyan Shi*, Kun Qian* (equal contribution), Xuwei Wang, Zhou Yu

Empirical Methods in Natural Language Processing (EMNLP) 2019

Persuasion for Good: Towards a Personalized Persuasive Dialogue System for Social Good **Best Paper Nomination**

Xuwei Wang*, Weiyan Shi* (equal contribution), Richard Kim, Yoojung Oh, Sijia Yang, Jingwen Zhang, Zhou Yu

Annual Meeting of the Association for Computational Linguistics (ACL) 2019

Unsupervised Dialogue Structure Learning

Weiyan Shi, Tiancheng Zhao, Zhou Yu

North American Chapter of the Association for Computational Linguistics (NAACL) 2019

Sentiment Adaptive End-to-End Dialogue Systems

Weiyan Shi, Zhou Yu

Annual Meeting of the Association for Computational Linguistics (ACL) 2018

Preprints

Just Fine-tune Twice: Selective Differential Privacy for Large Language Models

Weiyan Shi, Si Chen, Chiyuan Zhang, Ruoxi Jia, Zhou Yu

arXiv, 2022

Seamlessly Integrating Factual Information and Social Content with Persuasive Dialogue

Maximillian Chen, Weiyan Shi, Feifan Yan, Ryan Hou, Jingwen Zhang, Saurav Sahay, Zhou Yu

arXiv, 2022

GRANTS	Co-writer, “Selective Differential Privacy in Language Modeling”, Cisco Research Grant, 2022-2023. (PI: Zhou Yu)	
TALKS	<i>Reducing Repetition and Contradiction in Dialogue Systems</i> , Columbia NLP Seminar, Spring 2021 <i>Selective Differential Privacy in Language Modeling</i> , Columbia NLP Seminar, Fall 2021 <i>Sentiment-adaptive Dialogue Systems</i> , Cresta, Spring 2019	
MENTORING	Payton Guo Undergraduate student Project: Dialogue data denoising	2022-Present New York University
	Ryan Shea Master student Project: Membership inference attack on NLP models	2022-Present Columbia University
	Evan Li Undergraduate student Project: Differential privacy and attack models in NLP	2021 Columbia University
	Chelsea Cui Master student (Master at Columbia → Now at Amazon) Project: Differential privacy and dialogue simulation	2021 Columbia University
	Xuewei Wang Undergraduate student (Master at CMU → Now at Meta) Project: Persuasive dialogues and dialogue user simulator	2019-2020 Zhejiang University
TEACHING	Co-instructor: Conversational AI , Columbia University	Fall 2022
	Guest Lecturer: Natural Language Processing , Columbia University	Spring 2022

Topic: Dialogue System

Teaching Assistant: Conversational AI, Columbia University

Spring 2021

SERVICE

Area Chair: 17th Women in Machine Learning Workshop (WiML 2022)

Publicity Chair: 4th Workshop for Conversational AI, ACL 2022

Program Committee/Conference Reviewer:

ACL 2019, ACL 2020, *SEM 2020, ICLR 2021, AAAI 2021, EACL 2021, NAACL 2021, ACL 2021, *SEM 2021, NeurIPS 2021, EMNLP 2021, NLPCC 2021, ICLR 2022, AAAI 2022, ICML 2022, ACL 2022, CHI 2022, SIGDIAL 2022, HCI+NLP workshop at NAACL 2022, ACL Rolling Review 2021-2022

Journal Reviewer:

ACM Transactions on Human-Robot Interaction, Neurocomputing, ACM Transactions on Information Systems